

Feedback on the EchoPulse LinkedIn Playbook

TO: [Friend's name] **FROM:** Lakshya **RE:** The 20-archetype LinkedIn playbook + where I want to take it from here

Hey,

First — thank you. The 15-page playbook is the most useful single doc anyone has handed me on EchoPulse. I've been sitting with it for a few days, running it through how I'd actually operate week to week, and there's a lot in there that I want to keep verbatim. Before I get to the changes, I want to be specific about what's strong because the strong parts are *really* strong and I don't want them lost in the edits.

After that, I want to walk through what I'd change and — more importantly — *why*, because the changes only make sense if you understand the motives. I'm not editing for the sake of editing. Every adjustment maps to something specific I'm seeing about where EchoPulse is, who I'm selling to, and how LinkedIn (and Google) actually work in 2026. I want you to push back on any of these if you disagree.

At the end I'll show you the system I built around the playbook — three specialist agents that operationalize it — and tell you why I think pure-tactical playbooks always need that operating layer wrapped around them.

Let's go.

— Part 1 — What's strong (and what I'm keeping verbatim)

The 6 foundational rules

This is the single best part of the doc. The reason most LinkedIn writing advice is bad is that it stays at the level of "be authentic, post consistently, engage with comments." Your rules are operational and quantified:

- **Length: 700-1,500 chars; 20+ sentences = 1.18x engagement.** Specific. Testable. Defensible.

- **The hook is 90% of the post; line 1 max ~62 chars, line 2 max ~50 chars.** This is *the* rule that most playbooks miss. They tell you "have a strong hook" without telling you the dimensions of the actual hook. Yours has dimensions.
- **Short sentences, under 12 words. One thought per line. White space is your friend.** Yes. Non-negotiable.
- **End with a question. Reply in the first hour.** The first-hour rule alone is worth the whole doc — most people don't internalize how front-loaded LinkedIn distribution actually is.

I'm keeping all six rules unchanged in the revised version.

The 20-archetype catalog

Three things I love about how you structured these:

1. **Each archetype has *Why it works* + *Structure* + *EchoPulse draft*.** This three-part scaffold is what makes the catalog usable, not just inspirational. Most "post template" docs give you the template and stop there. Yours teaches the *system behind* each template, which means I can extrapolate to topics you didn't write a draft for.
2. **The drafts are voice-consistent across all 20.** That's hard. You clearly wrote them as one person, not 20 separate posts. The cadence carries.
3. **The mix of archetypes covers a real range — from "The Public Failure" (vulnerable) to "The Hot Take Listicle" (sharp) to "The Steal-This Resource" (utility).** This range is what prevents portfolio repetitiveness, which is the #1 reason audiences disengage at month 3.

The 20 archetypes are staying. I'm adding 5 more (which I'll explain in Part 3) but I'm not removing any.

The post-publish protocol

"First hour matters most. Engage on 5 other posts in your niche right after posting. Don't edit in the first hour. Track saves and DMs, not just likes."

These 4 rules are gold and most agencies don't follow them. Keeping verbatim.

The cadence framework

Mon Contrarian / Wed Utility / Fri Personal — that scaffolding gives the week a rhythm without locking it down. I'm keeping this exact structure as the default.

— Part 2 — What I'd change (and the motives behind each change)

Change #1 — Localize the drafts for US/CA/EU/UK buyers, not just Indian

Motive: EchoPulse's actual ICP is B2B founders and SaaS companies in the US, Canada, Western Europe, and the UK. The drafts in the playbook reference ₹6L, ₹2.5L, ₹40L, "Indian small businesses," "WhatsApp Business," "Bhopal chaiwala." These are wonderful for an Indian audience — they're *specific* and *authentic*, which is why they work. But for a US founder reading the post, "₹6L" requires conversion math, "WhatsApp Business" reads as a regional reference, and "Indian small businesses" filters them out of the audience.

This isn't an India-vs-US thing. It's a "where is the buyer" thing. Right now my buyer is in San Francisco, Toronto, Berlin, and Manchester. The drafts need to feel native to them.

The fix: Each archetype gets a parallel US/global draft alongside the Indian one. Same structure, same voice, different specifics: - "₹6L retainer" → "\$8K retainer" - "Indian small businesses" → "B2B founders" or "SaaS marketing leads" or "DTC brand owners" - "WhatsApp Business" → "email drip + LinkedIn DMs" - "Bhopal chaiwala" → keep when the post is about cross-domain analogies (the chaiwala story is a real Lakshya moment and works globally), but reference my actual lived experience, not made-up examples

Why this matters now, not later: The first 100 posts shape the audience LinkedIn assigns to me. If I post India-coded content for 90 days, LinkedIn shows my future posts to Indian audiences — and the Pilot pricing, retainer pricing, and growth motion are all calibrated to US/EU/UK budgets. The audience LinkedIn builds has compounding consequences.

Change #2 — Add bottom-of-funnel archetypes (5 new ones)

Motive: All 20 of your archetypes are top-of-funnel or middle-of-funnel — they earn audience, build authority, and create thought-leadership presence. None of them are designed for the moment a warm prospect is on the fence about working with us.

Right now, when a prospect is 80% sure they want to hire EchoPulse but waiting for one more signal, the playbook gives them more thought leadership. They don't need more thought leadership. They need: - A concrete preview of what working with us looks like - A clear filter for whether they're our kind of client - A pricing-and-rationale post that takes the awkwardness out of the buy decision

These are the posts that close.

The 5 new archetypes I'm adding:

#21 — The "Working With Us Looks Like" Post

A specific, non-pitchy walkthrough of what the first call, the kickoff doc, and week 2 deliverables actually feel like. Demystifies the buying decision.

#22 — The Anti-Positioning Post

"We don't work with X." A counter-intuitive filter post. Counter-intuitively builds demand because it signals strong opinion and high standards.

#23 — The Pricing Transparency Post

Full pricing, full rationale, full math. Most agencies hide pricing. Showing it earns trust *and* qualifies leads — the wrong-fit prospects self-disqualify, the right-fit ones lean in harder.

#24 — The Contrarian Case Study

A case study framed as a controversial choice. "We told a client to fire their best-performing channel" — same case study, but framed as POV, not as braggy results.

#25 — The Public Pricing Update

"We just dropped the Pilot from \$599 to \$299. Here's why." Pricing changes are inherently newsworthy and signal momentum. They earn distribution AND they qualify leads in the same post.

Why this matters now, not later: EchoPulse is 6 months old and just dropped Pilot pricing to \$299 to attract first 50 case studies. Every prospect who lands on my LinkedIn is going to ask, "Why so cheap? What's the catch? What do I actually get?" If my last 30 posts are all thought leadership, I can't answer those questions in-feed — I have to answer them in DM, which is slower and lower-trust. With bottom-of-funnel archetypes mixed in, the answer is already on my profile.

Change #3 — Replace "track saves and DMs" with a real measurement rubric

Motive: "Track saves and DMs, not just likes" is correct but not enough. It tells me *what* to track but not *what to do about what I tracked*. I need to know whether a specific post is a "wow keep doing this" or a "fine but skip in remix" or a "kill the topic, this isn't working." Without that signal, every post feels equally weighted and the engine drifts within a month.

The fix: A 7-tier post grading rubric (S / A / B / C / D / V / H), applied within 48 hours of every post. Each tier triggers a specific action: - **Tier S** (1000+ impressions, 50+ saves, 5+ DMs) → remix into 5+ derivatives, pin to profile - **Tier A** (500+ impressions, 25+ saves, 2+ DMs) → remix into 3 derivatives, cross-post to email - **Tier B** (200-500, 10-25 saves) → standard, move on - **Tier C** (under 200, under 10 saves) → diagnose, don't repeat the mistake - **Tier D** (under 50 impressions, under 5 saves) → kill the topic line if it happens 3

times in a row - **Bonus tier V** (viral) → 5,000+ impressions → mine the comments for next month's topics - **Bonus tier H** (high-intent) → low impressions, high-quality DMs → don't be fooled by impressions, this is the most valuable kind of post

The bonus H-tier is especially important because LinkedIn's vanity-metric trap punishes founders who optimize for likes. A post can do 80 impressions and produce 2 booked Pilots. That's a Tier H post. Most rubrics would call it a flop because the impressions are low. The real measurement is *did the right people see it*.

Why this matters now, not later: Without tiered grading, I'm flying blind on what to scale and what to kill. I'll keep posting the same archetypes that feel good to write but produce no DMs, and I won't notice the drift until I look back at quarter end.

Change #4 — Add a "Reply Lab" comment-engagement playbook

Motive: "Reply in the first hour" is right. But *how* to reply isn't in the doc. The reply itself is a mini-post — its quality determines whether the comment thread becomes a distribution event or stalls. Bad replies (one-line thanks, generic "great point") kill the conversation. Great replies (specific, additive, ending with a question back) extend the comment thread, which the algorithm reads as engagement signal and rewards with more distribution.

The fix: A reply playbook with templates for the 4 common comment types: 1.

Agreement ("Totally agree") → name the specific thing they agreed with + add a new detail not in the post + question back 2. **Disagreement** ("I think the opposite") → acknowledge the strongest part of their disagreement first + defend with specific example + concede one thing if true + invite elaboration 3. **Question** ("How would you handle X?") → answer directly in 2-3 sentences, don't deflect to DM unless answer is genuinely long 4.

Bot/template ("Great post ☐") → skip, don't engage

Plus tone guardrails: - Never use exclamation marks (calm > performative) - Never apologize for the opinion in a reply - Never say "thanks for engaging" — sounds like a chatbot - Never end with "Hope this helps!" — patronizing

Why this matters now, not later: First-hour comments are weighted 3x higher by LinkedIn's algorithm for distribution decisions. If I'm in the first hour of a hot post and replying generically, I'm leaving distribution on the table. Templated replies don't mean robotic — they mean having a tested structure ready so I'm not improvising under time pressure.

Change #5 — Add carousels and native video, not just text posts

Motive: The playbook is text-only. In 2026, LinkedIn's algorithm rewards multi-format diversity — accounts that mix text, carousel, document, native video, and live get more distribution per follower than text-only accounts. Text-only is a strategic ceiling.

This isn't about replacing text posts. Text posts are still the highest-leverage format for thought leadership and they're 60-70% of the portfolio. But: - **Carousels** are the highest-saved format on LinkedIn. Frameworks, step-by-step processes, comparisons all perform 2-3x better as carousels than as text. - **Native video (60-90 sec)** has the highest dwell time signal, which is one of LinkedIn's top algorithm inputs. - **Documents (PDFs)** earn massive reach from document opens — LinkedIn weighs that as deep engagement.

The fix: Each of the 25 archetypes gets a "format escalation" note — when does this archetype want a text post vs. a carousel vs. a native video? Some natural mappings: - The Cheat Sheet Drop (#7) → carousel preferred (visual list works better than text list) - The Steal-This Resource (#9) → document/PDF preferred (reader downloads the asset) - The Behind-the-Scenes Reality (#8) → native video preferred (humans on camera beats text) - The Contrarian Manifesto (#1) → text preferred (the language is the argument) - The Old vs. New (#14) → carousel preferred (visual contrast is the point)

Why this matters now, not later: Text-only accounts get capped on growth. By month 4 of text-only posting, the algorithm has decided what audience I have, and growth flattens. Mixing formats from the start re-trains LinkedIn weekly that this account is multi-modal and worth surfacing to new audiences.

Change #6 — Add a topic-mining + "Hill Test" filter

Motive: The playbook tells me *how* to write, but not *what* to write about. Topic selection is upstream of writing quality — most underperforming posts fail at topic, not at writing. A great writer can't save a bad topic.

The current implicit topic-selection method is "pick whatever feels right this week." That works for a week. It fails by month 3 — the audience can't articulate what EchoPulse stands for because the posts don't ladder up to a thesis.

The fix: A 6-source weekly topic sweep: 1. DMs received this week (highest signal — real audience asking real questions) 2. Sales call notes (highest signal — real objections, real "we tried X and it didn't work" stories) 3. Comments on previous posts (medium signal — what got people typing) 4. Competitor watchlist (medium signal — what's working in the niche) 5. Industry news / algo changes (medium signal — newsjack opportunities with 72h shelf life) 6. Search query data (medium signal — what is the audience Googling)

The convergence rule: if a topic shows up in 3+ sources in one week, write it. Period. That's the highest-signal pick available and it's almost never wrong.

Plus a 5-question Hill Test that every topic must pass before it gets calendared: 1. Will I defend this in a comment thread for an hour? 2. Is there a specific number, story, or example anchoring this? 3. Will the audience self-recognize in this post? 4. Does this differentiate EchoPulse, or repeat industry consensus? 5. Is there a measurable outcome I'd be proud of in 7 days?

A topic must pass all 5. If it doesn't, back to the backlog.

Why this matters now, not later: Without a topic-mining process, my Sunday-night content sessions become "what should I post about this week" panic. Random portfolios have low compounding because the audience can't recognize a pattern. With the 6-source sweep, every Monday I have 5-7 candidate topics, ranked by signal strength. The decision goes from agonizing to procedural.

Change #7 — Add a quarterly narrative + monthly themes

Motive: 50 posts in 6 months, picked individually, produce 50 disconnected impressions on the audience. 50 posts in 6 months that all ladder up to a quarterly narrative produce one cohesive impression: *I know exactly what this agency stands for.*

Without a narrative, individual posts can be great and the portfolio still feels random. With a narrative, even mediocre individual posts compound because each one reinforces the central position.

The fix: Every quarter has one narrative ("EchoPulse is the agency that lowers prices to earn its first 50 case studies"). Every month within the quarter has a theme that ladders up to it. Every week's thesis ladders up to the monthly theme.

Example for Q4 2026: - **Quarterly narrative:** "EchoPulse is the agency that drops Pilot prices to win first clients." - **Oct theme:** "The case for the \$299 Pilot" (transparency, why we cut, audience reactions) - **Nov theme:** "What we shipped in our first \$299 Pilots" (proof, behind-the-scenes, early case studies) - **Dec theme:** "What we learned from the first 30 Pilots" (synthesis, frameworks, year-end pieces)

Every Mon/Wed/Fri post in November ladders up to the November theme. Every November theme ladders up to the Q4 narrative. By December 31, the audience has had 90 days of one consistent message and can articulate what EchoPulse stands for in one sentence.

Why this matters now, not later: Position is what makes someone DM you. A founder doesn't DM the agency that produces good posts — they DM the agency that's clearly *for them*. The quarterly narrative is what makes that "clearly for them" feeling possible. Without it, I'm a generalist content agency that posts well, which competes with thousands.

Change #8 — Add a cross-format remix protocol

Motive: A great long-form blog deserves more than one publish. Once I've spent 3 hours on a 2,200-word piece, the marginal cost of producing 5 derivative formats is low (~30 min total) but the distribution multiplier is 5-7x. Most agencies skip the remix because they're already exhausted from writing the long-form. That's a leak.

The fix: Every Tier-A long-form artifact spawns 6 derivatives: 1. LinkedIn long-post (1,500 chars) — the single sharpest section, reframed 2. LinkedIn carousel (8-12 slides) — the framework visualized 3. Twitter / X thread — same framework, atomized 4. Email to list — the blog summary + one "behind-the-scenes" take that's NOT in the blog (rewards subscribers) 5. YouTube short / Reel script (60 sec) — hook + 3 framework points + CTA 6. Two future-post seeds — sub-topics the blog couldn't fully cover

Done weekly, this turns 3 hours of writing into 6 distribution surfaces. Compounding.

Why this matters now, not later: I'm one person. I can't outproduce a team. I can outdistribute one by making each piece work 6x harder. The remix protocol is the leverage that solo founders need to compete with agencies that have dedicated content teams.

— Part 3 — The system I built around the playbook

Here's the mental model shift I want to share with you: **the playbook is a tactical reference. It's not an operating system.** A solo founder running a content engine can't execute from a tactical reference alone — they need an operating system that decides *what* to write before they sit down to write it, and grades *what worked* after publishing.

I built that operating system as three specialist agents:

1. **The Viral Post Writer** — Owns LinkedIn distribution. Takes topics from the Marketing Head, runs them through the Hook Lab, ships 3 posts a week with first-hour reply pre-writes. Doesn't decide what to write — executes brilliantly on what's been decided.
2. **The Value Content Creator** — Owns evergreen + Google index. Takes the week's anchor topic from the Marketing Head, produces a 1,500-2,500 word blog with SEO discipline, and designs it for cross-format remix into 5-6 derivatives.
3. **The Marketing Head / Analyst** — Owns the brain. Picks topics via the 6-source sweep, runs the Hill Test, sets the weekly thesis, monitors competitors, grades every post on the 7-tier rubric, decides what to scale and what to kill, and orchestrates the other two.

The key insight: **the Marketing Head decides; the writers execute.** The most common failure mode for solo content operations is the same brain doing strategy + writing + measurement — and all three suffer because of context-switching. Splitting the

roles (even when one person plays all three at different times of the week) makes the operation cleaner.

Total time investment: ~4-5 hours per week. Output: 3 LinkedIn posts + 1 long-form blog + 1 carousel + 5+ days of comment engagement + a measurement report.

I have full operational specs written for each of the three agents (one ~2000-word doc per agent). They include workflows, quality gates, output templates, anti-patterns, and worked examples. If you want to read them, I can share — they're under [marketing-agents/](#) in my workspace folder.

— Part 4 — Suggested next steps

Three things I'd love your take on:

1. **Do these motives track for you?** Specifically: am I right that EchoPulse needs US-localized drafts, bottom-of-funnel archetypes, and a measurement rubric? Or do you think the original playbook is closer to right than I'm giving it credit for?
2. **What did I miss?** You've watched LinkedIn evolve longer than I have. Is there a category of post or a measurement signal that the original playbook *and* my additions both skip?
3. **Want to co-author the v2?** I can take a first stab at the 5 new bottom-of-funnel archetypes (drafts + Why It Works + Structure, in your original format) and you can red-pen them. Or you take a first stab and I red-pen. Either way, I want the v2 to feel like ours, not just mine.

The original playbook is genuinely the best LinkedIn doc I've read for an agency. The system I built around it isn't a replacement — it's the operating layer that makes the catalog ship-able week after week without my brain melting. I want to credit the catalog as the foundation everywhere I use it.

Hit me back when you've had a chance to read this. Coffee on me when we sync.

— Lakshya